Technical Critical Review

Disentangled Speaker Representations via Adversarial Autoencoder with GRL and MAPC Regularisation

Saumya Pancholi M25CSA027 IIT Jodhpur

1 Problem Statement

Speaker verification systems must decide whether two utterances originate from the same speaker. A fundamental obstacle is that embeddings extracted from raw audio conflate two statistically independent sources of variation: *speaker identity* (vocal tract shape, prosody) and *environmental factors* (channel response, room acoustics, additive noise). A model that cannot separate these will produce embeddings that shift with recording condition, directly inflating the equal error rate (EER) when enrolment and test conditions differ.

The reviewed method poses this as a *disentanglement* problem: learn a factorised representation $\mathbf{e} \to (\mathbf{e}_S, \mathbf{e}_E)$ such that $\mathbf{e}_S$ captures only speaker-discriminative information and $\mathbf{e}_E$ captures only environment-specific variation, while retaining all information needed to reconstruct the original embedding. This is a well-motivated problem, as cross-condition robustness remains one of the main failure modes of deployed speaker verification systems.

2 Proposed Method

The method builds on ECAPA-TDNN [?] — a strong backbone using SE-Res2 blocks and attentive statistics pooling — and appends three complementary disentanglement modules:

2.1 Disentangling Autoencoder (AE)

Two parallel linear encoders split the backbone embedding into speaker and environment subspaces: $f_S : \mathbf{e} \to \mathbf{e}_S \in \mathbb{R}^{d_S}$ and $f_E : \mathbf{e} \to \mathbf{e}_E \in \mathbb{R}^{d_E}$. A shared decoder reconstructs $\hat{\mathbf{e}} = g([\mathbf{e}_S; \mathbf{e}_E])$, and the reconstruction loss $\mathcal{L}_r = \|\mathbf{e} - \hat{\mathbf{e}}\|_1$ ensures no information is discarded.

2.2 Adversarial Environment Suppression (GRL)

A gradient-reversal layer [?] is attached to $\mathbf{e}_S$, feeding an environment discriminator. During backpropagation, gradients are negated by factor α, forcing the encoder f_S to produce representations from which environment cannot be predicted. The adversarial loss $\mathcal{L}_{\text{adv}}$ uses a triplet formulation on discriminator outputs.

2.3 MAPC Decorrelation

A maximum a-posteriori correlation penalty minimises linear statistical dependence between subspaces:

$$\mathcal{L}_{\text{MAPC}} = \frac{1}{B} \left\| \tilde{\mathbf{e}}_S^\top \tilde{\mathbf{e}}_E \right\|_1 \tag{1}$$

where $\tilde{\cdot}$ denotes column-normalised, mean-centred activations.

2.4 Speaker Classification (AAM-Softmax)

Additive angular margin softmax on $\mathbf{e}_S$ enforces speaker-discriminability. The combined objective is:

$$\mathcal{L} = \lambda_{\text{spk}}\mathcal{L}_{\text{AAM}} + \lambda_r\mathcal{L}_r + \lambda_{ee}\mathcal{L}_{\text{triplet}} + \lambda_{\text{adv}}\mathcal{L}_{\text{GRL}} + \lambda_c\mathcal{L}_{\text{MAPC}} \tag{2}$$

At inference, only $\mathbf{e}_S$ is used for cosine similarity scoring, making the approach a drop-in replacement for standard embedding extraction.

3 Strengths

Orthogonal regularisation. GRL operates on gradients during backpropagation while MAPC operates on activations directly. This dual-pressure design means environment information is penalised at two independent points in the computation graph, reducing the likelihood that either constraint is individually bypassed by the optimiser.

Information-preserving split. The reconstruction loss prevents the degenerate solution where $\mathbf{e}_S \to \mathbf{0}$, which would trivially satisfy the adversarial objective at the cost of all speaker information. This is a non-obvious but important design choice.

Strong backbone. ECAPA-TDNN is a proven state-of-the-art backbone for speaker verification. Applying disentanglement at the embedding level (rather than modifying the feature extractor) means the full representational power of the backbone is available to both subspaces.

No domain labels required. The method learns environment invariance purely from triplet augmentation, making it applicable without explicit recording-condition metadata, which is rarely available in practice.

4 Weaknesses

Linear disentanglement. The AE encoders are single linear layers. Speaker and environment factors interact non-linearly — room acoustics modulate formant amplitudes in a frequency-dependent manner — which a linear separator cannot model. Non-linear encoders (MLPs or convolutional blocks) would be more expressive, at the cost of added optimisation complexity.

Implicit environment assumption. The method assumes that data augmentation (noise addition, reverb) provides sufficient implicit "environment labels". If augmented pairs share channel characteristics (e.g., all data is from the same microphone), $\mathbf{e}_E$ may inadvertently absorb speaker-correlated spectral variation rather than genuine environmental variation.

Hyperparameter sensitivity. Five loss weights must be jointly tuned. The GRL scale α requires a careful warmup schedule to prevent gradient explosion during early training. The original paper provides limited ablation on these values, making reproducibility from the paper alone challenging.

Subspace dimension choice. The split $d_S + d_E$ is fixed as a hyperparameter. There is no mechanism to learn the intrinsic dimensionality of each factor, meaning the model may over-allocate capacity to one subspace.

5 Assumptions

The method rests on four key assumptions: (1) speaker identity and environment are *statistically independent* sources of variation, which is not strictly true (different speakers may appear predominantly in certain acoustic environments); (2) a *linear* projection is sufficient to achieve the separation; (3) *triplet augmentation* provides representative environment contrast; and (4) the speaker verification task is well-served by a *cosine similarity* metric on the disentangled $\mathbf{e}_S$, implying that post-disentanglement embeddings lie on a hypersphere with uniform angular density per speaker.

6 Experimental Validity

Table 1: Summary of experimental concerns.

Concern	**Detail**
Single dataset	Evaluated on one corpus; no cross-dataset generalisation test (e.g., VoxCeleb $\rightarrow$ SITW)
No confidence int.	Point estimates only; no bootstrap CI or significance test across seeds
No ablation table	Individual contributions of GRL, MAPC, AE not isolated
Short training	3-epoch budget prevents convergence; val-EER still decreasing at epoch 3
Closed-set eval	Test speakers overlap with training distribution; open-set generalisation untested

The absence of a component-wise ablation study is the most critical gap. Without it, it is impossible to determine whether the EER improvement is driven by GRL, MAPC, the reconstruction loss, or simply by the dimension reduction from $\mathbf{e}$ to $\mathbf{e}_S$. A convincing experimental section would require: (i) ablation of each loss term, (ii) evaluation on a held-out cross-condition test set, and (iii) multiple runs with different random seeds to establish variance estimates.

Despite these limitations, the direction is well-motivated, and the observed reduction in EER from 23.17% (baseline) to 10.83% (disentangled) on the reproduction experiment provides empirical evidence that the disentanglement provides genuine benefit under the tested conditions.